

# ChurnPrediction

September 10, 2024

## 1 Welcome to the Data Science Coding Challange!

Test your skills in a real-world coding challenge. Coding Challenges provide CS & DS Coding Competitions with Prizes and achievement badges!

CS & DS learners want to be challenged as a way to evaluate if they're job ready. So, why not create fun challenges and give winners something truly valuable such as complimentary access to select Data Science courses, or the ability to receive an achievement badge on their Coursera Skills Profile - highlighting their performance to recruiters.

### 1.1 Introduction

In this challenge, you'll get the opportunity to tackle one of the most industry-relevant machine learning problems with a unique dataset that will put your modeling skills to the test. Subscription services are leveraged by companies across many industries, from fitness to video streaming to retail. One of the primary objectives of companies with subscription services is to decrease churn and ensure that users are retained as subscribers. In order to do this efficiently and systematically, many companies employ machine learning to predict which users are at the highest risk of churn, so that proper interventions can be effectively deployed to the right audience.

In this challenge, we will be tackling the churn prediction problem on a very unique and interesting group of subscribers on a video streaming service!

Imagine that you are a new data scientist at this video streaming company and you are tasked with building a model that can predict which existing subscribers will continue their subscriptions for another month. We have provided a dataset that is a sample of subscriptions that were initiated in 2021, all snapshotted at a particular date before the subscription was cancelled. Subscription cancellation can happen for a multitude of reasons, including:

- \* the customer completes all content they were interested in, and no longer need the subscription
- \* the customer finds themselves to be too busy and cancels their subscription until a later time
- \* the customer determines that the streaming service is not the best fit for them, so they cancel and look for something better suited

Regardless the reason, this video streaming company has a vested interest in understanding the likelihood of each individual customer to churn in their subscription so that resources can be allocated appropriately to support customers. In this challenge, you will use your machine learning toolkit to do just that!

## 1.2 Understanding the Datasets

### 1.2.1 Train vs. Test

In this competition, you'll gain access to two datasets that are samples of past subscriptions of a video streaming platform that contain information about the customer, the customers streaming preferences, and their activity in the subscription thus far. One dataset is titled `train.csv` and the other is titled `test.csv`.

`train.csv` contains 70% of the overall sample (243,787 subscriptions to be exact) and importantly, will reveal whether or not the subscription was continued into the next month (the "ground truth").

The `test.csv` dataset contains the exact same information about the remaining segment of the overall sample (104,480 subscriptions to be exact), but does not disclose the "ground truth" for each subscription. It's your job to predict this outcome!

Using the patterns you find in the `train.csv` data, predict whether the subscriptions in `test.csv` will be continued for another month, or not.

### 1.2.2 Dataset descriptions

Both `train.csv` and `test.csv` contain one row for each unique subscription. For each subscription, a single observation (`CustomerID`) is included during which the subscription was active.

In addition to this identifier column, the `train.csv` dataset also contains the target label for the task, a binary column `Churn`.

Besides that column, both datasets have an identical set of features that can be used to train your model to make predictions. Below you can see descriptions of each feature. Familiarize yourself with them so that you can harness them most effectively for this machine learning task!

```
[1]: import pandas as pd
data_descriptions = pd.read_csv('data_descriptions.csv')
pd.set_option('display.max_colwidth', None)
data_descriptions
```

Column_name	Column_type	Data_type	Description
0 AccountAge	Feature	integer	The age of the user's account in months.
1 MonthlyCharges	Feature	float	The amount charged to the user on a monthly basis.
2 TotalCharges	Feature	float	The total charges incurred by the user over the account's lifetime.
3 SubscriptionType	Feature	object	The type of subscription chosen by the user (Basic, Standard, or Premium).
4 PaymentMethod	Feature	string	The method of payment used by the user.
5 PaperlessBilling	Feature	string	Indicates whether the user has opted for paperless billing (Yes or No).

6	ContentType	Feature	string	The type of content →preferred by the user (Movies, TV Shows, or Both).
7	MultiDeviceAccess	Feature	string	Indicates whether the user →has access to the service on multiple devices (Yes or No).
8	DeviceRegistered	Feature	string	The type of device →registered by the user (TV, Mobile, Tablet, or Computer).
9	ViewingHoursPerWeek	Feature	float	The number of hours the user →spends watching content per week.
10	AverageViewingDuration	Feature	float	The average duration of →each viewing session in minutes.
11	ContentDownloadsPerMonth	Feature	integer	The number of content →downloads by the user per month.
12	GenrePreference	Feature	string	The preferred genre of →content chosen by the user.
13	UserRating	Feature	float	The user's rating for the →service on a scale of 1 to 5.
14	SupportTicketsPerMonth	Feature	integer	The number of support →tickets raised by the user per month.
15	Gender	Feature	string	The gender of the user (Male →or Female).
16	WatchlistSize	Feature	float	The number of items in the →user's watchlist.
17	ParentalControl	Feature	string	Indicates whether parental →control is enabled for the user (Yes or No).
18	SubtitlesEnabled	Feature	string	Indicates whether subtitles →are enabled for the user (Yes or No).
19	CustomerID	Identifier	string	A unique identifier for each →customer.
20	Churn	Target	integer	The target variable →indicating whether a user has churned or not (1 for churned, 0 for not →churned).

### 1.3 How to Submit your Predictions to Coursera

Submission Format:

In this notebook you should follow the steps below to explore the data, train a model using the data in `train.csv`, and then score your model using the data in `test.csv`. Your final submission should be a dataframe (call it `prediction_df` with two columns and exactly 104,480 rows (plus a header row). The first column should be `CustomerID` so that we know which prediction belongs to which observation. The second column should be called `predicted_probability` and should be a numeric column representing the **likelihood that the subscription will churn**.

Your submission will show an error if you have extra columns (beyond `CustomerID` and `predicted_probability`) or extra rows. The order of the rows does not matter.

The naming convention of the dataframe and columns are critical for our autograding, so please make sure to use the exact naming conventions of `prediction_df` with column names `CustomerID`

and predicted\_probability!

To determine your final score, we will compare your predicted\_probability predictions to the source of truth labels for the observations in test.csv and calculate the [ROC AUC](#). We choose this metric because we not only want to be able to predict which subscriptions will be retained, but also want a well-calibrated likelihood score that can be used to target interventions and support most accurately.

## 1.4 Import Python Modules

First, import the primary modules that will be used in this project. Remember as this is an open-ended project please feel free to make use of any of your favorite libraries that you feel may be useful for this challenge. For example some of the following popular packages may be useful:

- pandas
- numpy
- Scipy
- Scikit-learn
- keras
- matplotlib
- seaborn

```
[2]: # Import required packages

# Data packages
import pandas as pd
import numpy as np

# Machine Learning / Classification packages
from sklearn.metrics import roc_auc_score
from sklearn.model_selection import train_test_split
from sklearn.dummy import DummyClassifier

from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_auc_score

# Visualization Packages
from matplotlib import pyplot as plt
import seaborn as sns
%matplotlib inline
```

## 1.5 Load the Data

Let's start by loading the dataset `train.csv` into a dataframe `train_df`, and `test.csv` into a dataframe `test_df` and display the shape of the dataframes.

```
[3]: train_df = pd.read_csv("train.csv")
print('train_df Shape:', train_df.shape)
train_df.head()
```

train\_df Shape: (243787, 21)

```
[3]:   AccountAge  MonthlyCharges  TotalCharges  SubscriptionType \
0          20        11.055215    221.104302      Premium
1          57        5.175208     294.986882      Basic
2          73        12.106657    883.785952      Basic
3          32        7.263743     232.439774      Basic
4          57        16.953078    966.325422      Premium

      PaymentMethod PaperlessBilling ContentType MultiDeviceAccess \
0    Mailed check            No       Both           No
1    Credit card            Yes      Movies           No
2    Mailed check            Yes      Movies           No
3  Electronic check          No    TV Shows           No
4  Electronic check          Yes    TV Shows           No

  DeviceRegistered ViewingHoursPerWeek  ...  ContentDownloadsPerMonth \
0        Mobile             36.758104  ...                  10
1        Tablet             32.450568  ...                  18
2      Computer             7.395160  ...                  23
3        Tablet             27.960389  ...                  30
4         TV                20.083397  ...                  20

  GenrePreference UserRating  SupportTicketsPerMonth  Gender WatchlistSize \
0      Sci-Fi      2.176498                      4  Male            3
1      Action      3.478632                      8  Male           23
2    Fantasy      4.238824                      6  Male            1
3      Drama      4.276013                      2  Male           24
4    Comedy      3.616170                      4 Female           0

  ParentalControl SubtitlesEnabled CustomerID Churn
0            No            No  CB6SXPNVZA    0
1            No            Yes  S7R2G87009    0
2           Yes            Yes  EASDC20BDT    0
3           Yes            Yes  NPF69NT69N    0
4            No            No  4LGYPK7VOL    0
```

[5 rows x 21 columns]

```
[4]: test_df = pd.read_csv("test.csv")
print('test_df Shape:', test_df.shape)
test_df.head()
```

test\_df Shape: (104480, 20)

```
[4]:   AccountAge  MonthlyCharges  TotalCharges  SubscriptionType \
0          38        17.869374     679.036195      Premium
1          77        9.912854      763.289768      Basic
2           5       15.019011      75.095057    Standard
3          88       15.357406     1351.451692    Standard
4          91       12.406033     1128.949004    Standard

      PaymentMethod PaperlessBilling ContentType MultiDeviceAccess \
0      Mailed check            No    TV Shows            No
1  Electronic check           Yes    TV Shows            No
2    Bank transfer            No    TV Shows           Yes
3  Electronic check           No      Both            Yes
4    Credit card             Yes    TV Shows           Yes

  DeviceRegistered ViewingHoursPerWeek AverageViewingDuration \
0            TV           29.126308         122.274031
1            TV           36.873729         57.093319
2          Computer          7.601729        140.414001
3           Tablet          35.586430        177.002419
4           Tablet          23.503651         70.308376

  ContentDownloadsPerMonth GenrePreference UserRating \
0                  42        Comedy     3.522724
1                  43        Action     2.021545
2                  14        Sci-Fi     4.806126
3                  14        Comedy     4.943900
4                   6        Drama     2.846880

  SupportTicketsPerMonth  Gender WatchlistSize ParentalControl \
0                      2    Male        23            No
1                      2  Female        22           Yes
2                      2  Female        22            No
3                      0  Female        23           Yes
4                      6  Female         0            No

  SubtitlesEnabled CustomerID
0            No  O1W6BHP6RM
1            No  LFR4X92X8H
2           Yes  QM5GBIYODA
3           Yes  D9RXTK2K9F
4            No  ENTCCHR1LR
```

## 1.6 Explore, Clean, Validate, and Visualize the Data (optional)

Feel free to explore, clean, validate, and visualize the data however you see fit for this competition to help determine or optimize your predictive model. Please note - the final autograding will only be on the accuracy of the prediction\_df predictions.


```
[5]: # Remove duplicates if any
train_df.drop_duplicates(inplace=True)

# Handle missing values: for example, filling missing numeric values with mean
train_df.fillna(train_df.mean(), inplace=True)
```

```
[6]: # Check the distribution of the target variable
print(train_df['Churn'].value_counts(normalize=True))


# Boxplot for numeric features
numeric_columns = train_df.select_dtypes(include=['float64', 'int64']).columns
plt.figure(figsize=(15, 8))
sns.boxplot(data=train_df[numeric_columns])
plt.xticks(rotation=90)
plt.show()
```

```
0    0.818768
1    0.181232
Name: Churn, dtype: float64
```




```
[7]: # Heatmap to see correlations between features
```


```
plt.figure(figsize=(10, 8))
sns.heatmap(train_df.corr(), annot=True, cmap='coolwarm')
plt.title('Correlation Heatmap')
plt.show()
```




```
[8]: # Visualize the distribution of numeric columns
train_df[numeric_columns].hist(bins=30, figsize=(15, 10))
plt.show()
```




```
[9]: # Boxplot to visualize how MonthlyCharges are distributed by Churn
plt.figure(figsize=(10, 6))
sns.boxplot(x='Churn', y='MonthlyCharges', data=train_df)
plt.title('Monthly Charges vs Churn')
plt.show()
```



```
[10]: # Boxplot to visualize how ViewingHoursPerWeek are distributed by Churn
plt.figure(figsize=(10, 6))
sns.boxplot(x='Churn', y='ViewingHoursPerWeek', data=train_df)
plt.title('Viewing Hours Per Week vs Churn')
plt.show()
```



```
[11]: # Check distribution of categorical columns (e.g., SubscriptionType) based on Churn
plt.figure(figsize=(12, 6))
sns.countplot(x='SubscriptionType', hue='Churn', data=train_df)
plt.title('Subscription Type vs Churn')
plt.show()
```



## 1.7 Make predictions (required)

Remember you should create a dataframe named prediction\_df with exactly 104,480 entries plus a header row attempting to predict the likelihood of churn for subscriptions in test\_df. Your submission will throw an error if you have extra columns (beyond CustomerID and predicted\_probaility) or extra rows.

The file should have exactly 2 columns: CustomerID (sorted in any order) predicted\_probability (contains your numeric predicted probabilities between 0 and 1, e.g. from estimator.predict\_proba(X, y)[:, 1])

The naming convention of the dataframe and columns are critical for our autograding, so please make sure to use the exact naming conventions of prediction\_df with column names CustomerID and predicted\_probability!

### 1.7.1 Example prediction submission:

The code below is a very naive prediction method that simply predicts churn using a Dummy Classifier. This is used as just an example showing the submission format required. Please change/alter/delete this code below and create your own improved prediction methods for generating prediction\_df.

**PLEASE CHANGE CODE BELOW TO IMPLEMENT YOUR OWN PREDICTIONS**

```
[12]: # Define features and target
X = train_df.drop(columns=['Churn', 'CustomerID'])
y = train_df['Churn']
X_test = test_df.drop(columns=['CustomerID'])
customer_ids = test_df['CustomerID']

[13]: # Preprocess data
numeric_features = ['AccountAge', 'MonthlyCharges', 'TotalCharges', ↳
    'ViewingHoursPerWeek', 'AverageViewingDuration', 'ContentDownloadsPerMonth', ↳
    'UserRating', 'SupportTicketsPerMonth', 'WatchlistSize']
categorical_features = ['SubscriptionType', 'PaymentMethod', 'PaperlessBilling', ↳
    'ContentType', 'MultiDeviceAccess', 'DeviceRegistered', 'GenrePreference', ↳
    'Gender', 'ParentalControl', 'SubtitlesEnabled']

# Create preprocessing pipelines for numeric and categorical features
numeric_transformer = StandardScaler()
categorical_transformer = OneHotEncoder(handle_unknown='ignore')

# Combine preprocessing steps
preprocessor = ColumnTransformer(
    transformers=[('num', numeric_transformer, numeric_features),
```

```

        ('cat', categorical_transformer, categorical_features)
    ])

```

[14]: # Create and train the model pipeline  
model = Pipeline(steps=[('preprocessor', preprocessor),  
 ('classifier', LogisticRegression(max\_iter=1000))])

[15]: # Split the data  
X\_train, X\_valid, y\_train, y\_valid = train\_test\_split(X, y, test\_size=0.2, random\_state=42)  
  
# Train the model  
model.fit(X\_train, y\_train)

[15]: Pipeline(memory=None,  
 steps=[('preprocessor',  
 ColumnTransformer(n\_jobs=None, remainder='drop',  
 sparse\_threshold=0.3,  
 transformer\_weights=None,  
 transformers=[('num',  
 StandardScaler(copy=True,  
 with\_mean=True,  
 with\_std=True),  
 ['AccountAge',  
 'MonthlyCharges',  
 'TotalCharges',  
 'ViewingHoursPerWeek',  
 'AverageViewingDuration',  
 'ContentDownloadsPerMonth',  
 'UserRating',  
 'Sup...',  
 'GenrePreference', 'Gender',  
 'ParentalControl',  
 'SubtitlesEnabled'])],  
 verbose=False)),  
 ('classifier',  
 LogisticRegression(C=1.0, class\_weight=None, dual=False,  
 fit\_intercept=True, intercept\_scaling=1,  
 l1\_ratio=None, max\_iter=1000,  
 multi\_class='auto', n\_jobs=None,  
 penalty='l2', random\_state=None,  
 solver='lbfgs', tol=0.0001, verbose=0,  
 warm\_start=False)],  
 verbose=False)

[18]: # Predict probabilities on the test data

```

X_test_scaled = model.named_steps['preprocessor'].transform(X_test) # Use the preprocessor to transform test data
predicted_probabilities = model.predict_proba(X_test)[:, 1]

# Create the prediction dataframe
prediction_df = pd.DataFrame({
    'CustomerID': customer_ids,
    'predicted_probability': predicted_probabilities
})

```

[20]:

```

# Validate the model
y_valid_probs = model.predict_proba(X_valid)[:, 1]
# Calculate the ROC AUC score
roc_auc = roc_auc_score(y_valid, y_valid_probs)
print(f'ROC AUC Score: {roc_auc}')

```

ROC AUC Score: 0.7536603799535646

[21]:

```

print(prediction_df.shape)
print(prediction_df.head(10))

```

	CustomerID	predicted_probability
0	01W6BHP6RM	0.105559
1	LFR4X92X8H	0.041280
2	QM5GBIYODA	0.399395
3	D9RXTK2K9F	0.044968
4	ENTCCHR1LR	0.135402
5	7A88BB5I06	0.470202
6	700MW9XEWR	0.118827
7	EL1RMFMPYL	0.291825
8	4IA2QPT6ZK	0.204686
9	AEDCWHSJDN	0.182445

## 1.8 Final Tests - IMPORTANT - the cells below must be run prior to submission

Below are some tests to ensure your submission is in the correct format for autograding. The autograding process accepts a csv prediction\_submission.csv which we will generate from our prediction\_df below. Please run the tests below and ensure no assertion errors are thrown.

```
[22]: # FINAL TEST CELLS - please make sure all of your code is above these test cells
```

```
# Writing to csv for autograding purposes
prediction_df.to_csv("prediction_submission.csv", index=False)
submission = pd.read_csv("prediction_submission.csv")

assert isinstance(submission, pd.DataFrame), 'You should have a dataframe named  
→prediction_df.'
```

```
[23]: # FINAL TEST CELLS - please make sure all of your code is above these test cells
```

```
assert submission.columns[0] == 'CustomerID', 'The first column name should be  
→CustomerID.'
assert submission.columns[1] == 'predicted_probability', 'The second column name  
→should be predicted_probability.'
```

```
[24]: # FINAL TEST CELLS - please make sure all of your code is above these test cells
```

```
assert submission.shape[0] == 104480, 'The dataframe prediction_df should have  
→104480 rows.'
```

```
[25]: # FINAL TEST CELLS - please make sure all of your code is above these test cells
```

```
assert submission.shape[1] == 2, 'The dataframe prediction_df should have 2  
→columns.'
```

## 1.9 SUBMIT YOUR WORK!

Once we are happy with our prediction\_df and prediction\_submission.csv we can now submit for autograding! Submit by using the blue **Submit Assignment** at the top of your notebook. Don't worry if your initial submission isn't perfect as you have multiple submission attempts and will obtain some feedback after each submission!