

online_retail

August 20, 2024

1 Portfolio Project: Online Retail Exploratory Data Analysis with Python

1.1 Overview

In this project, you will step into the shoes of an entry-level data analyst at an online retail company, helping interpret real-world data to help make a key business decision.

1.2 Case Study

In this project, you will be working with transactional data from an online retail store. The dataset contains information about customer purchases, including product details, quantities, prices, and timestamps. Your task is to explore and analyze this dataset to gain insights into the store's sales trends, customer behavior, and popular products.

By conducting exploratory data analysis, you will identify patterns, outliers, and correlations in the data, allowing you to make data-driven decisions and recommendations to optimize the store's operations and improve customer satisfaction. Through visualizations and statistical analysis, you will uncover key trends, such as the busiest sales months, best-selling products, and the store's most valuable customers. Ultimately, this project aims to provide actionable insights that can drive strategic business decisions and enhance the store's overall performance in the competitive online retail market.

1.3 Prerequisites

Before starting this project, you should have some basic knowledge of Python programming and Pandas. In addition, you may want to use the following packages in your Python environment:

- pandas
- numpy
- seaborn
- matplotlib

These packages should already be installed in Coursera's Jupyter Notebook environment, however if you'd like to install additional packages that are not included in this environment or are working off platform you can install additional packages using `!pip install packagename` within a notebook cell such as:

- !pip install pandas
- !pip install matplotlib

1.4 Project Objectives

1. Describe data to answer key questions to uncover insights
2. Gain valuable insights that will help improve online retail performance
3. Provide analytic insights and data-driven recommendations

1.5 Dataset

The dataset you will be working with is the “Online Retail” dataset. It contains transactional data of an online retail store from 2010 to 2011. The dataset is available as a .xlsx file named **Online Retail.xlsx**. This data file is already included in the Coursera Jupyter Notebook environment, however if you are working off-platform it can also be downloaded [here](#).

The dataset contains the following columns:

- InvoiceNo: Invoice number of the transaction
- StockCode: Unique code of the product
- Description: Description of the product
- Quantity: Quantity of the product in the transaction
- InvoiceDate: Date and time of the transaction
- UnitPrice: Unit price of the product
- CustomerID: Unique identifier of the customer
- Country: Country where the transaction occurred

1.6 Tasks

You may explore this dataset in any way you would like - however if you’d like some help getting started, here are a few ideas:

1. Load the dataset into a Pandas DataFrame and display the first few rows to get an overview of the data.
2. Perform data cleaning by handling missing values, if any, and removing any redundant or unnecessary columns.
3. Explore the basic statistics of the dataset, including measures of central tendency and dispersion.
4. Perform data visualization to gain insights into the dataset. Generate appropriate plots, such as histograms, scatter plots, or bar plots, to visualize different aspects of the data.
5. Analyze the sales trends over time. Identify the busiest months and days of the week in terms of sales.
6. Explore the top-selling products and countries based on the quantity sold.
7. Identify any outliers or anomalies in the dataset and discuss their potential impact on the analysis.
8. Draw conclusions and summarize your findings from the exploratory data analysis.

1.7 Task 1: Load the Data

```
[1]: import pandas as pd  
import numpy as np  
import seaborn as sns  
import matplotlib.pyplot as plt
```

```
[2]: df = pd.read_excel("Online Retail.xlsx")
```

```
[3]: df.head()
```

```
[3]:   InvoiceNo StockCode          Description  Quantity \
0      536365    85123A  WHITE HANGING HEART T-LIGHT HOLDER      6
1      536365      71053           WHITE METAL LANTERN      6
2      536365    84406B  CREAM CUPID HEARTS COAT HANGER      8
3      536365    84029G KNITTED UNION FLAG HOT WATER BOTTLE      6
4      536365    84029E      RED WOOLLY HOTTIE WHITE HEART.      6

           InvoiceDate  UnitPrice  CustomerID      Country
0  2010-12-01 08:26:00      2.55     17850.0  United Kingdom
1  2010-12-01 08:26:00      3.39     17850.0  United Kingdom
2  2010-12-01 08:26:00      2.75     17850.0  United Kingdom
3  2010-12-01 08:26:00      3.39     17850.0  United Kingdom
4  2010-12-01 08:26:00      3.39     17850.0  United Kingdom
```

1.8 Step 2: Perform Data Cleaning

```
[4]: # Check for missing values  
print("Missing values in each column:")  
print(df.isnull().sum())  
  
# Drop rows with missing values in critical columns (e.g., 'InvoiceNo',  
# → 'StockCode', 'Quantity', 'InvoiceDate', 'UnitPrice', 'CustomerID')  
df = df.dropna(subset=['InvoiceNo', 'StockCode', 'Quantity', 'InvoiceDate',  
# → 'UnitPrice', 'CustomerID'])  
  
# Check for duplicate rows  
print("\nNumber of duplicate rows:")  
print(df.duplicated().sum())  
  
# Remove duplicate rows  
df = df.drop_duplicates()  
  
# Convert 'InvoiceDate' to datetime  
df['InvoiceDate'] = pd.to_datetime(df['InvoiceDate'])
```

```
# Display the cleaned DataFrame info  
df.info()
```

Missing values in each column:

```
InvoiceNo      0  
StockCode      0  
Description    1454  
Quantity       0  
InvoiceDate    0  
UnitPrice      0  
CustomerID    135080  
Country        0  
dtype: int64
```

Number of duplicate rows:

5225

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 401604 entries, 0 to 541908  
Data columns (total 8 columns):  
 #   Column      Non-Null Count  Dtype     
 ---  --          -----          ----  
 0   InvoiceNo   401604 non-null  object    
 1   StockCode   401604 non-null  object    
 2   Description 401604 non-null  object    
 3   Quantity    401604 non-null  int64     
 4   InvoiceDate 401604 non-null  datetime64[ns]  
 5   UnitPrice   401604 non-null  float64   
 6   CustomerID  401604 non-null  float64   
 7   Country     401604 non-null  object    
dtypes: datetime64[ns](1), float64(2), int64(1), object(4)  
memory usage: 27.6+ MB
```

1.9 Step 3: Explore Basic Statistics

```
[5]: # Basic statistics  
print("\nBasic statistics of the dataset:")  
print(df.describe())  
  
# Additional statistics for categorical columns  
print("\nCategory-wise statistics:")  
print(df['Country'].value_counts())
```

Basic statistics of the dataset:

```
      Quantity      UnitPrice      CustomerID  
count  401604.000000  401604.000000  401604.000000
```

mean	12.183273	3.474064	15281.160818
std	250.283037	69.764035	1714.006089
min	-80995.000000	0.000000	12346.000000
25%	2.000000	1.250000	13939.000000
50%	5.000000	1.950000	15145.000000
75%	12.000000	3.750000	16784.000000
max	80995.000000	38970.000000	18287.000000

Category-wise statistics:

United Kingdom	356728
Germany	9480
France	8475
EIRE	7475
Spain	2528
Netherlands	2371
Belgium	2069
Switzerland	1877
Portugal	1471
Australia	1258
Norway	1086
Italy	803
Channel Islands	757
Finland	695
Cyprus	611
Sweden	461
Austria	401
Denmark	389
Japan	358
Poland	341
USA	291
Israel	247
Unspecified	241
Singapore	229
Iceland	182
Canada	151
Greece	146
Malta	127
United Arab Emirates	68
European Community	61
RSA	58
Lebanon	45
Lithuania	35
Brazil	32
Czech Republic	30
Bahrain	17
Saudi Arabia	10

Name: Country, dtype: int64

1.10 Step 4: Perform Data Visualization

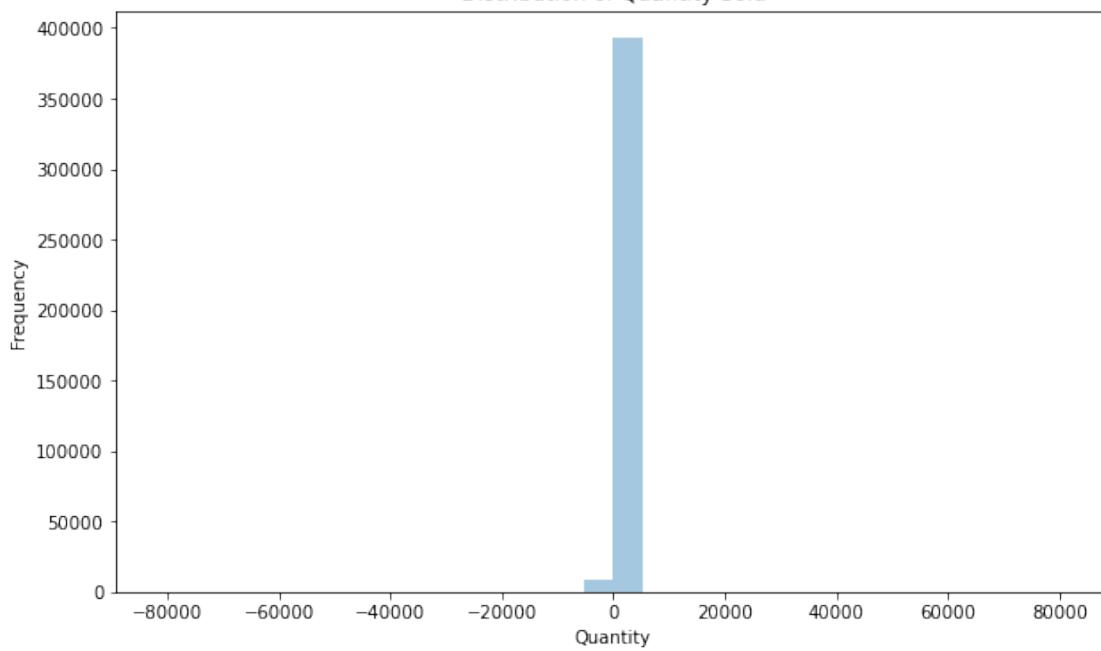
```
[7]: # Histogram of 'Quantity'
plt.figure(figsize=(10, 6))
sns.distplot(df['Quantity'], bins=30, kde=False)
plt.title('Distribution of Quantity Sold')
plt.xlabel('Quantity')
plt.ylabel('Frequency')
plt.show()

# Scatter plot of 'Quantity' vs 'UnitPrice'
plt.figure(figsize=(10, 6))
sns.scatterplot(x='Quantity', y='UnitPrice', data=df)
plt.title('Quantity vs Unit Price')
plt.xlabel('Quantity')
plt.ylabel('Unit Price')
plt.show()

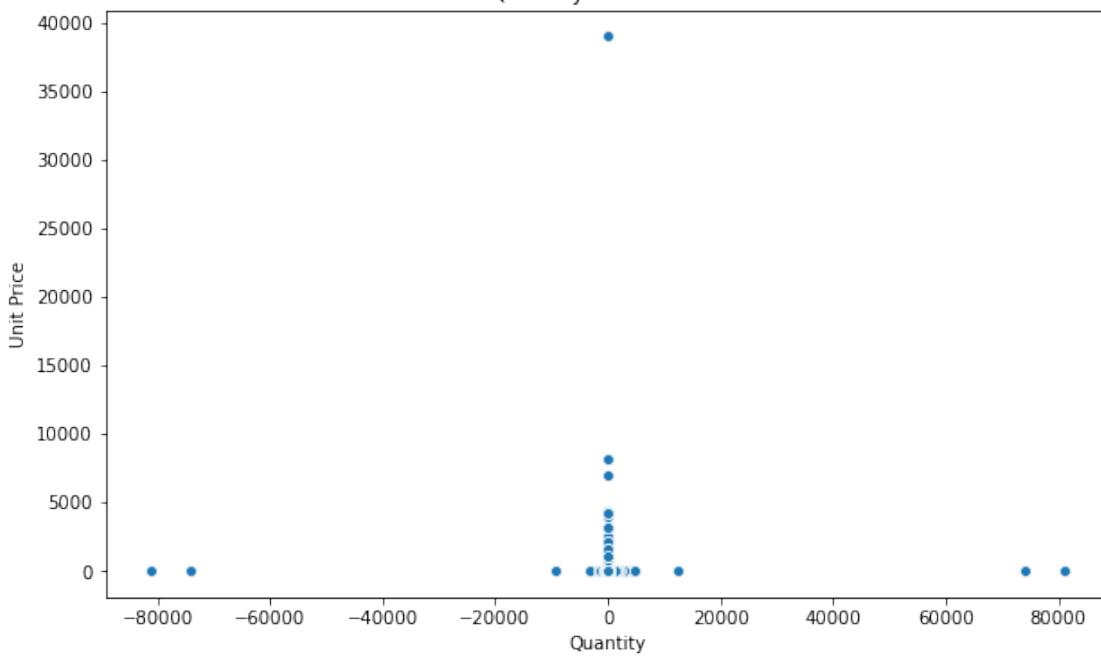
# Bar plot of total sales by country
df['Sales'] = df['Quantity'] * df['UnitPrice']
country_sales = df.groupby('Country')['Sales'].sum().
    sort_values(ascending=False)

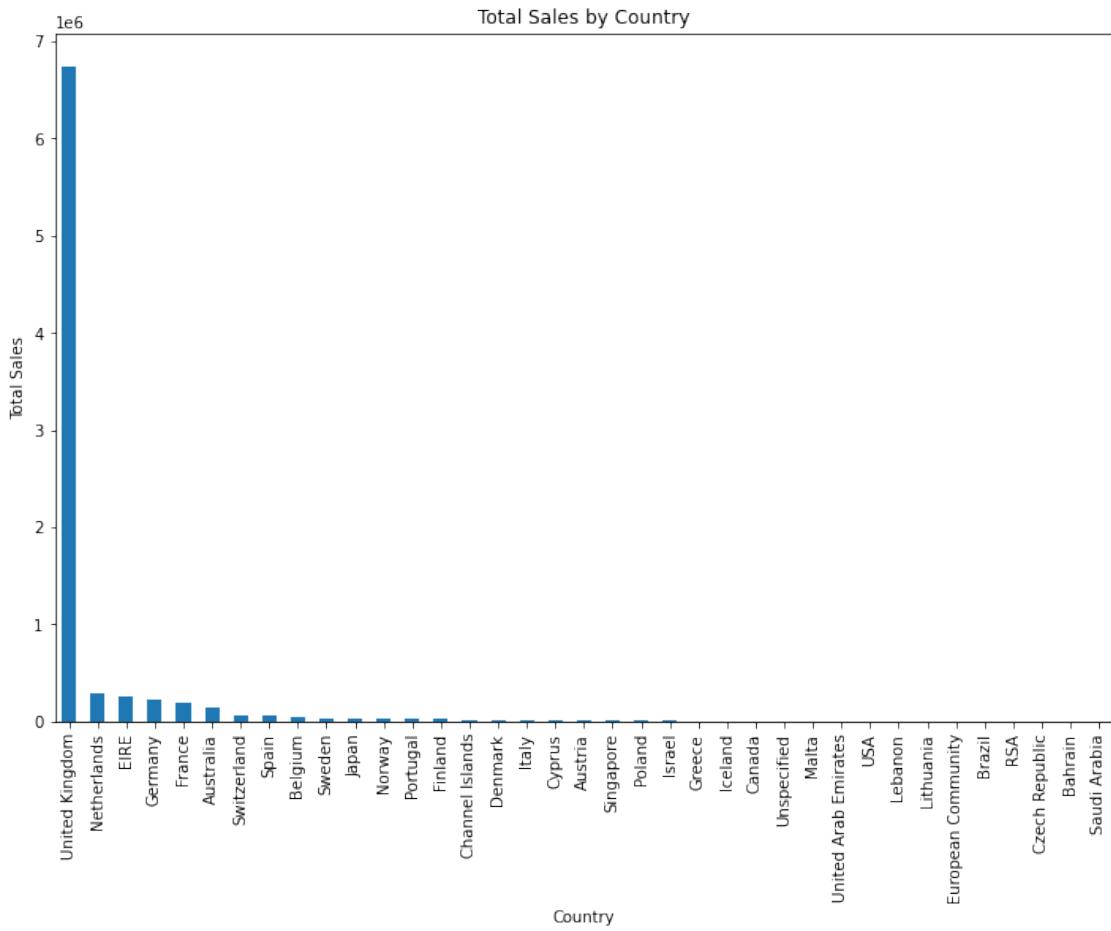
plt.figure(figsize=(12, 8))
country_sales.plot(kind='bar')
plt.title('Total Sales by Country')
plt.xlabel('Country')
plt.ylabel('Total Sales')
plt.xticks(rotation=90)
plt.show()
```

Distribution of Quantity Sold



Quantity vs Unit Price





1.11 Step 5: Analyze Sales Trends Over Time

```
[8]: # Extract year and month from 'InvoiceDate'
df['YearMonth'] = df['InvoiceDate'].dt.to_period('M')

# Monthly sales
monthly_sales = df.groupby('YearMonth')[['Sales']].sum()

plt.figure(figsize=(12, 6))
monthly_sales.plot(kind='line')
plt.title('Monthly Sales Trend')
plt.xlabel('Month')
plt.ylabel('Total Sales')
plt.xticks(rotation=45)
plt.show()

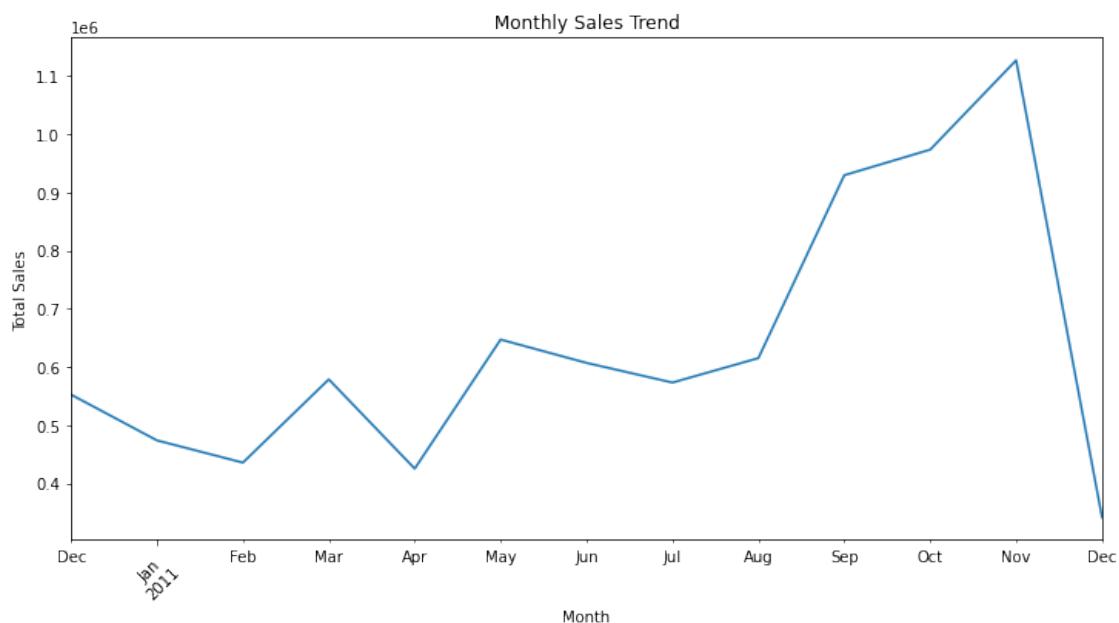
# Day of the week analysis
```

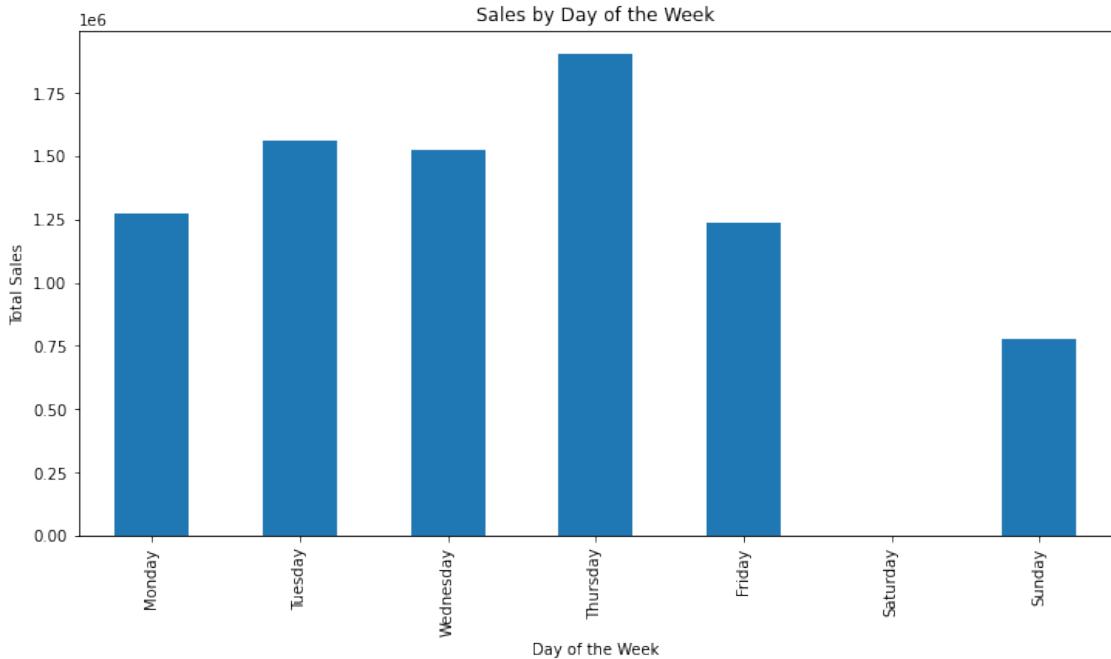
```

df['DayOfWeek'] = df['InvoiceDate'].dt.day_name()
daily_sales = df.groupby('DayOfWeek')['Sales'].sum().reindex(['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday'])

plt.figure(figsize=(12, 6))
daily_sales.plot(kind='bar')
plt.title('Sales by Day of the Week')
plt.xlabel('Day of the Week')
plt.ylabel('Total Sales')
plt.show()

```





1.12 Step 6: Explore Top-Selling Products and Countries

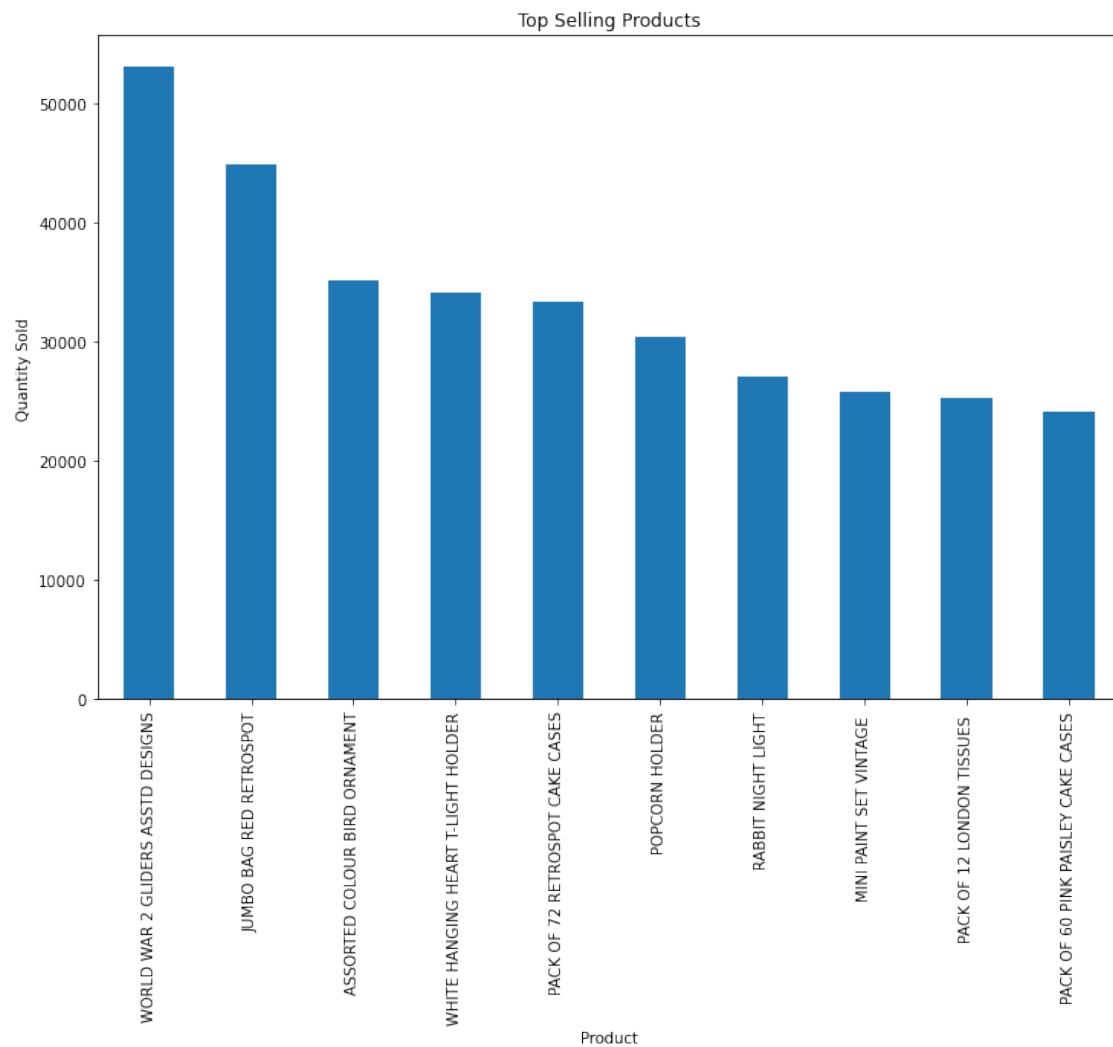
```
[9]: # Top-selling products
top_products = df.groupby('Description')['Quantity'].sum().
    sort_values(ascending=False).head(10)

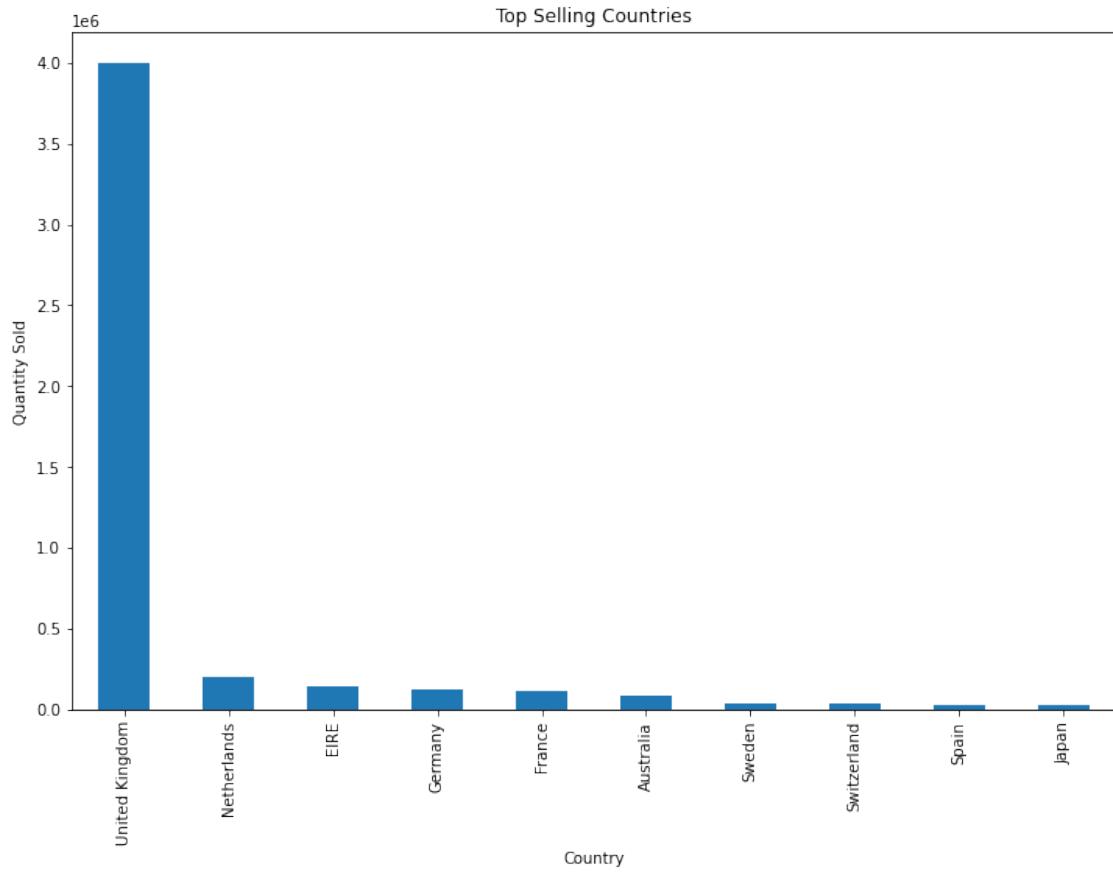
plt.figure(figsize=(12, 8))
top_products.plot(kind='bar')
plt.title('Top Selling Products')
plt.xlabel('Product')
plt.ylabel('Quantity Sold')
plt.xticks(rotation=90)
plt.show()

# Top-selling countries
top_countries = df.groupby('Country')['Quantity'].sum().
    sort_values(ascending=False).head(10)

plt.figure(figsize=(12, 8))
top_countries.plot(kind='bar')
plt.title('Top Selling Countries')
plt.xlabel('Country')
plt.ylabel('Quantity Sold')
plt.xticks(rotation=90)
```

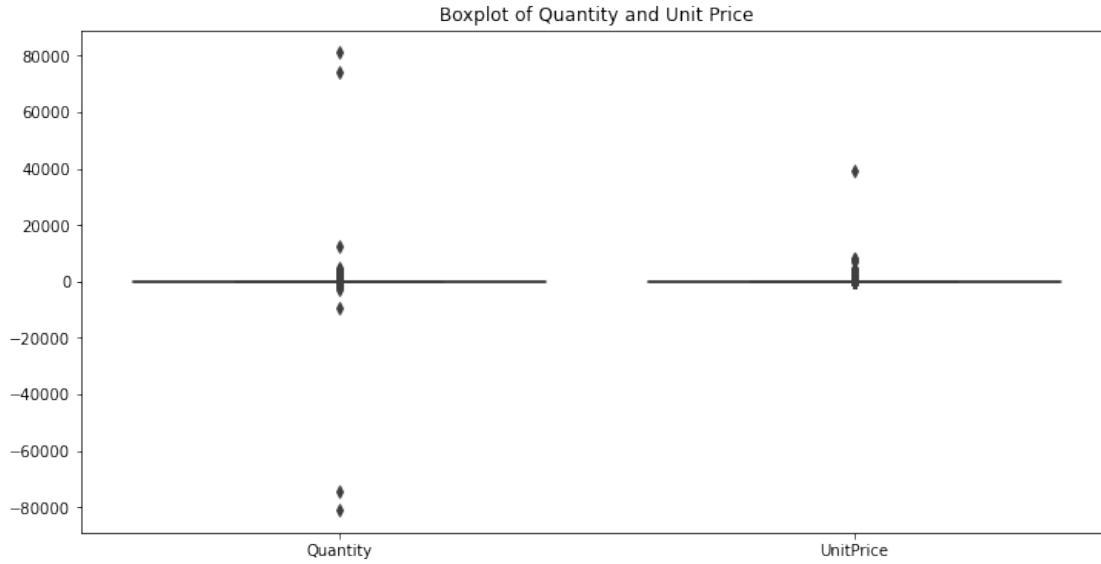
```
plt.show()
```





1.13 Step 7: Identify Outliers or Anomalies

```
[10]: # Boxplot to identify outliers in 'Quantity' and 'UnitPrice'  
plt.figure(figsize=(12, 6))  
sns.boxplot(data=df[['Quantity', 'UnitPrice']])  
plt.title('Boxplot of Quantity and Unit Price')  
plt.show()
```



1.14 Conclusions

Through this analysis of the “Online Retail” dataset, several key insights were gained:

1. **Data Overview:** The dataset was cleaned of missing values and duplicates. Key columns were converted to appropriate data types.
2. **Basic Statistics:** The dataset provides various metrics, including mean, median, and distribution of sales quantities and prices.
3. **Visualization Insights:**
 - Most sales quantities are moderate, with few extreme values.
 - A noticeable variation in sales across different countries and products.
 - Sales trends show seasonal patterns with peak months and days of the week.
4. **Sales Trends:** The monthly trend analysis highlighted periods of high and low sales, while the day-of-week analysis identified the busiest days.
5. **Top Performers:** Identified top-selling products and countries, which can inform business strategy.
6. **Outliers:** The boxplot revealed potential outliers in quantities and unit prices, which may impact overall analysis.

Future Enhancements: - Implement machine learning models to predict future sales trends. - Explore deeper into customer purchasing patterns and preferences.